AutoML for Neural Network Robustness Verification

Matthias König | ADA Workshop on AutoAI

09.12.2022

Universiteit
Leiden
The Netherlands

Neural networks are vulnerable to adversarial examples

Neural networks are vulnerable to adversarial examples

Some examples… and possible consequences…

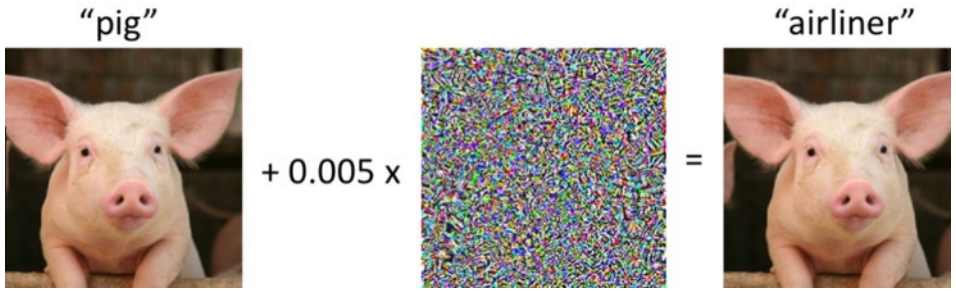

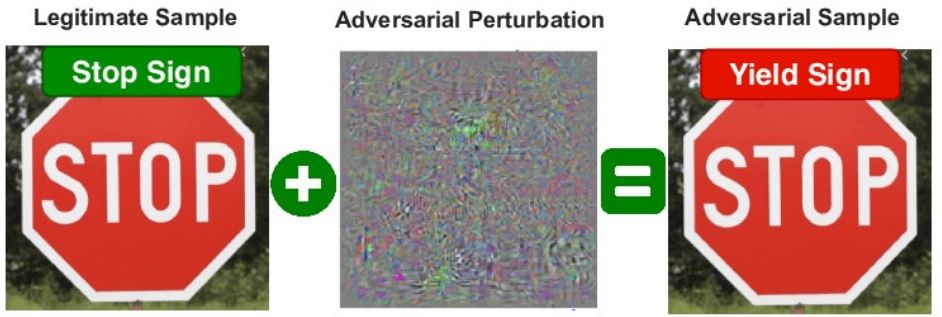

Verifying a deep neural network

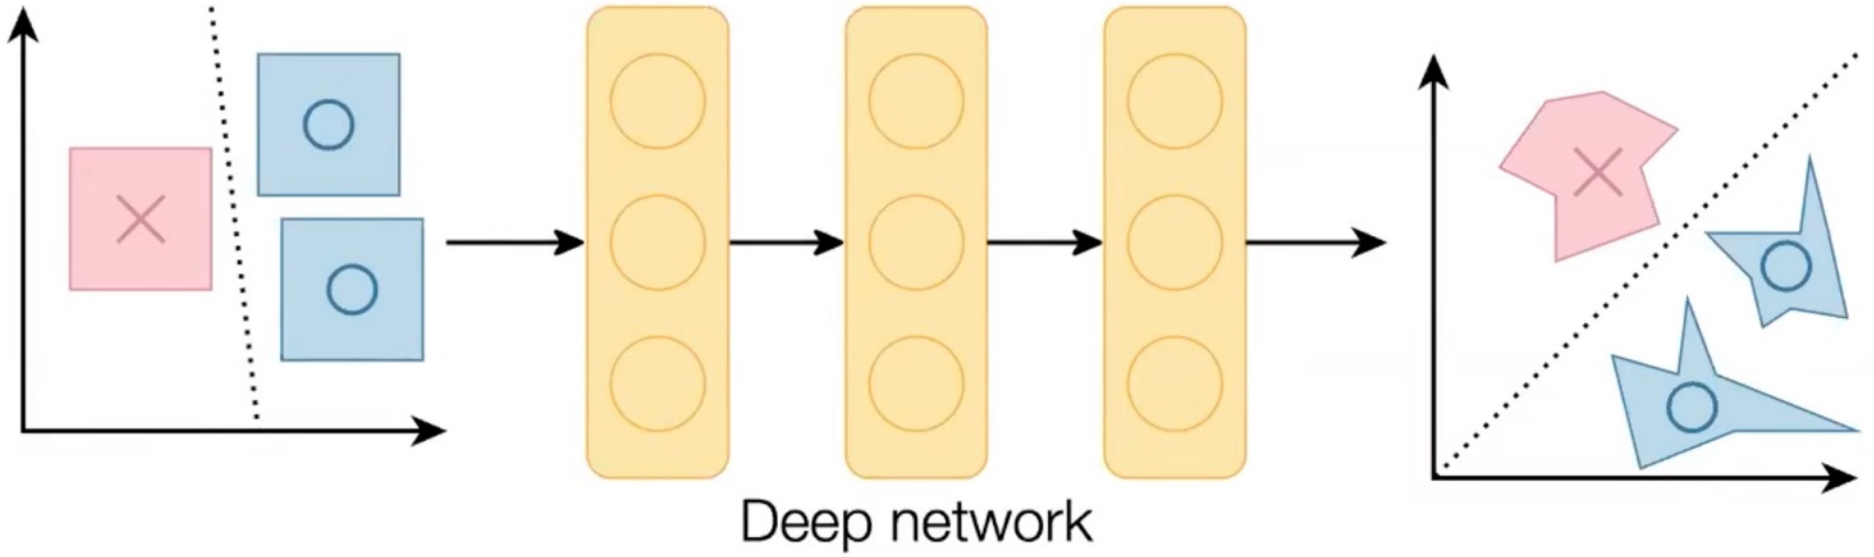

Image source: Stanford AI Safety Seminar

Neural network verification can be expensive

Incomplete vs. complete verification

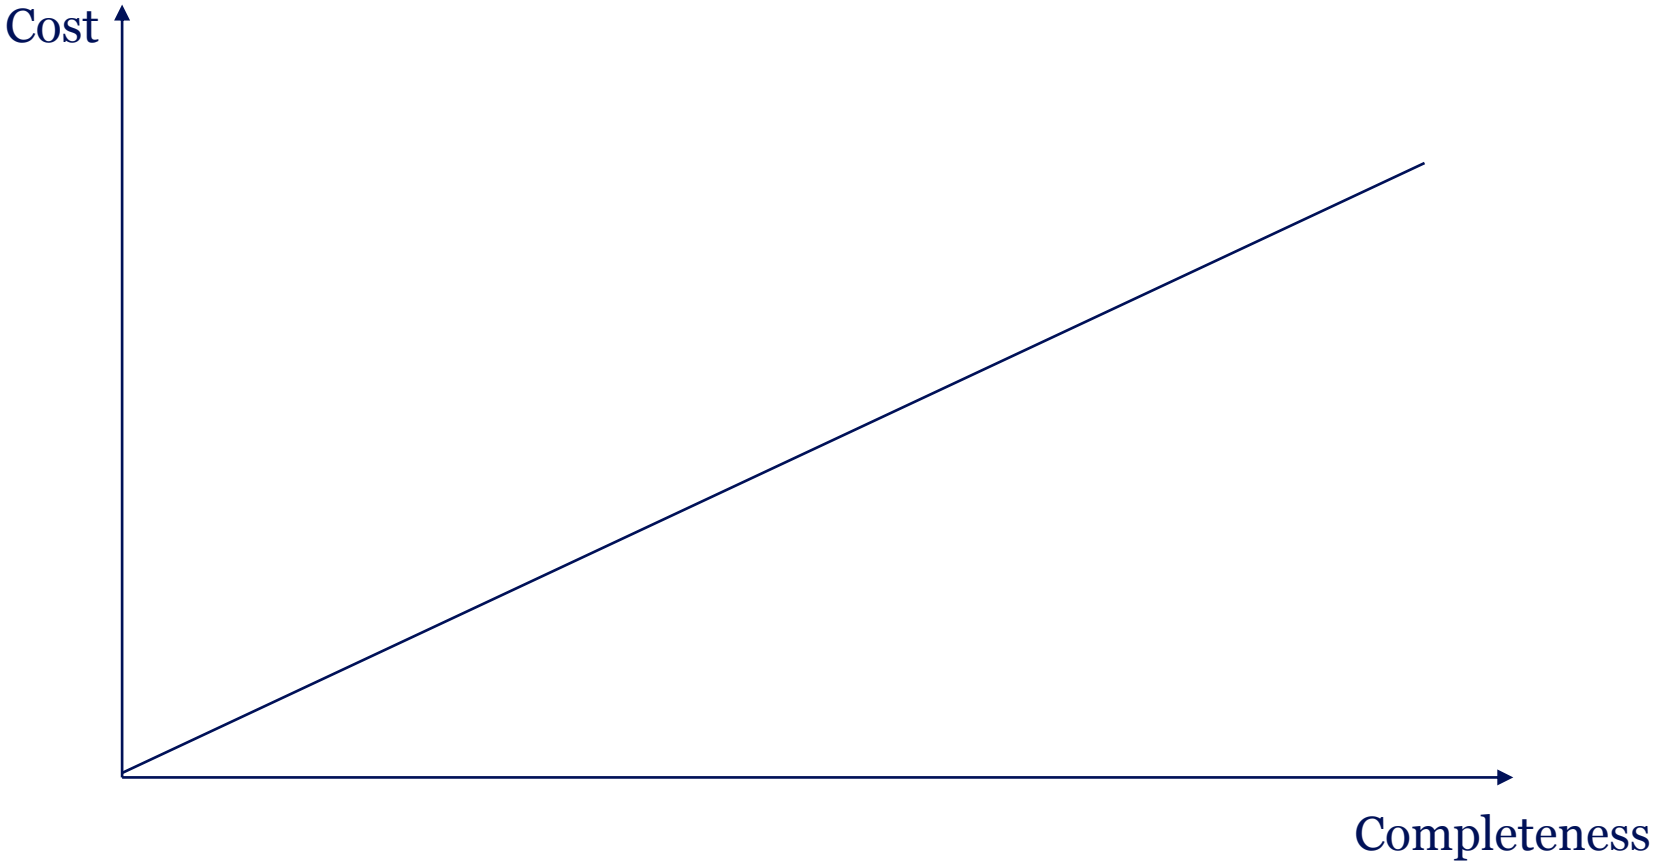

Neural network verification can be expensive

Incomplete vs. complete verification

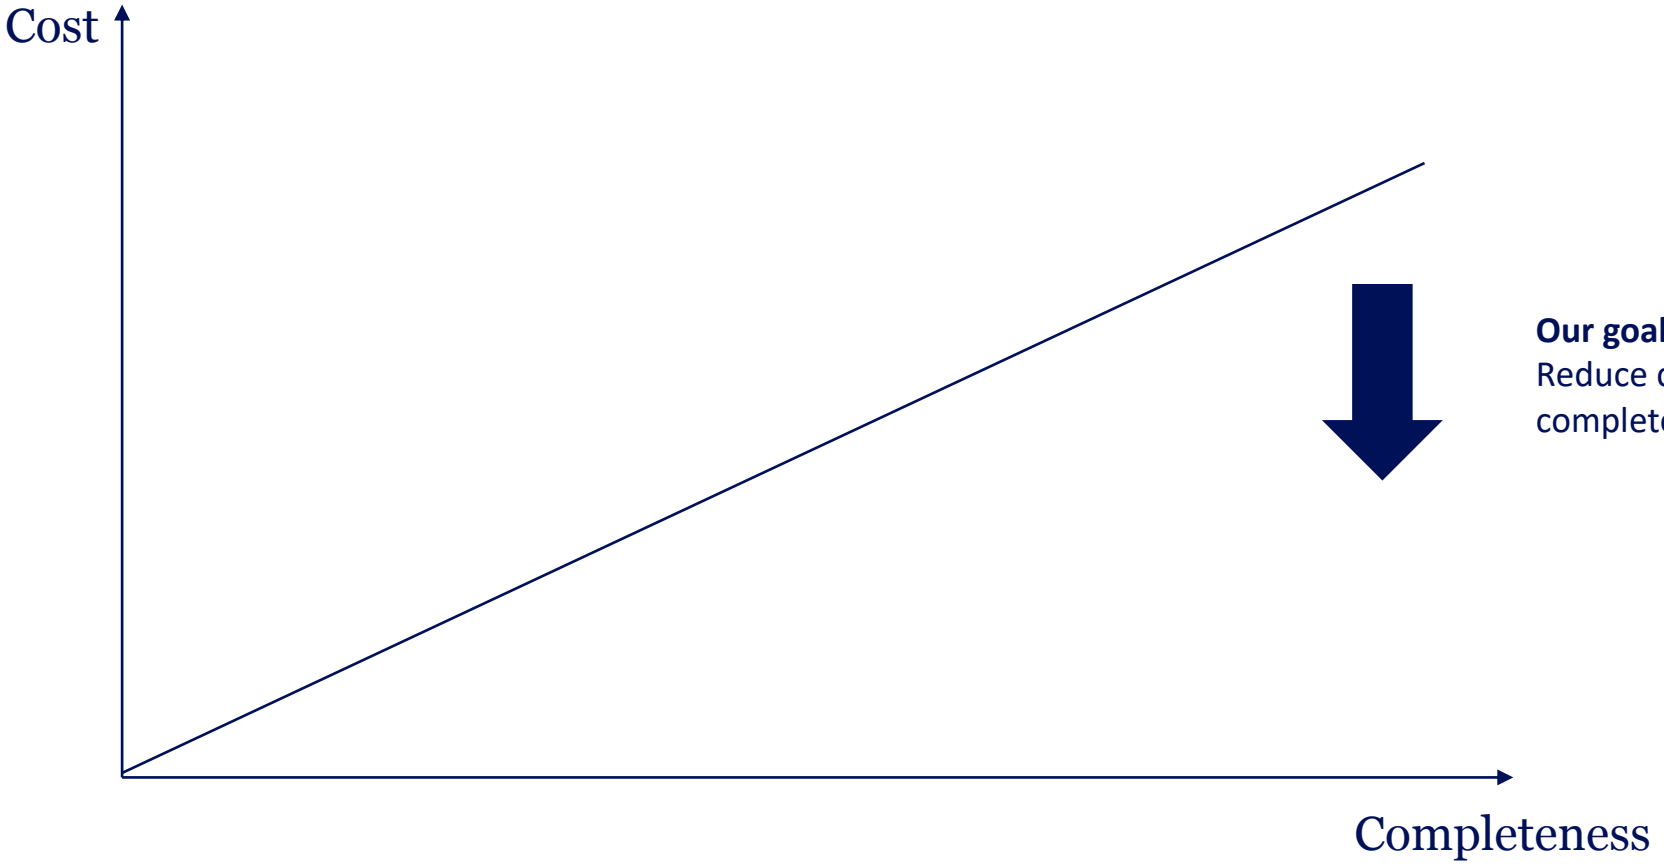

Our goal:
Reduce computational cost of complete verification methods

There exist several approaches to verify a network

Some examples…

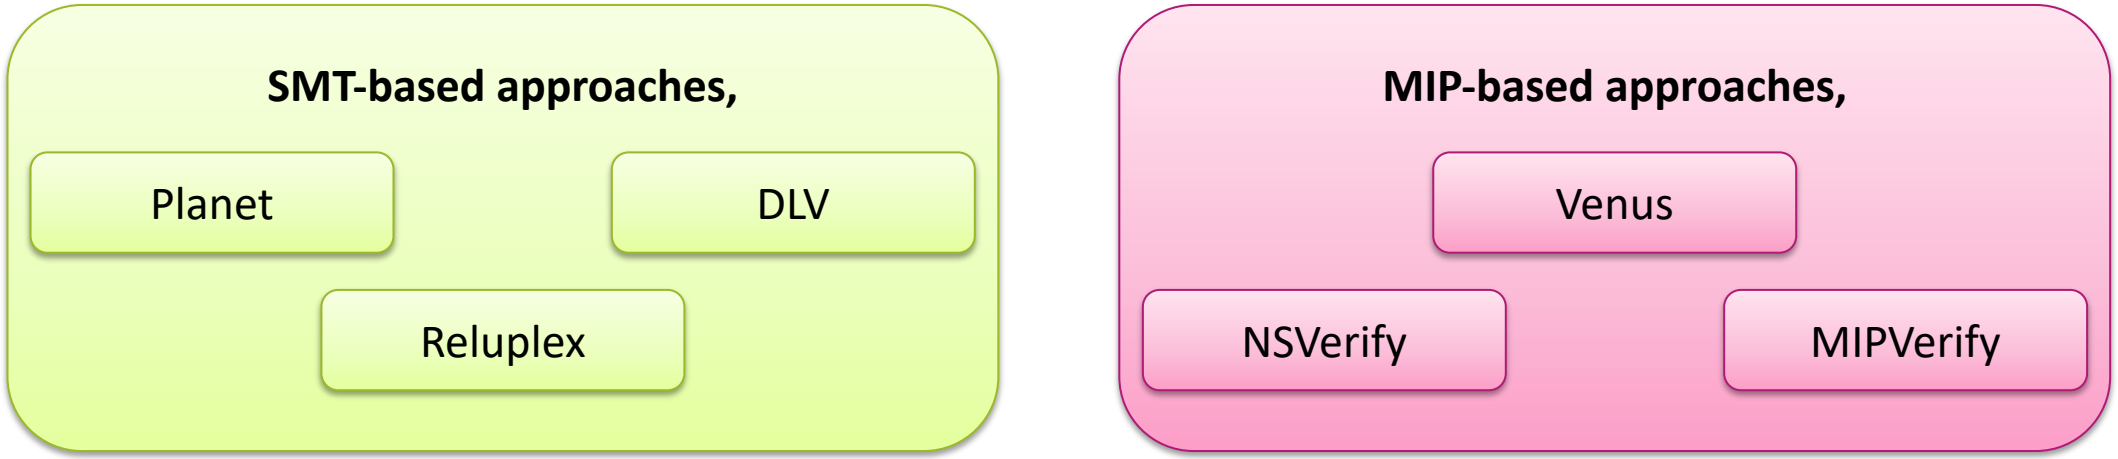

There exist several approaches to verify a network

Some examples…

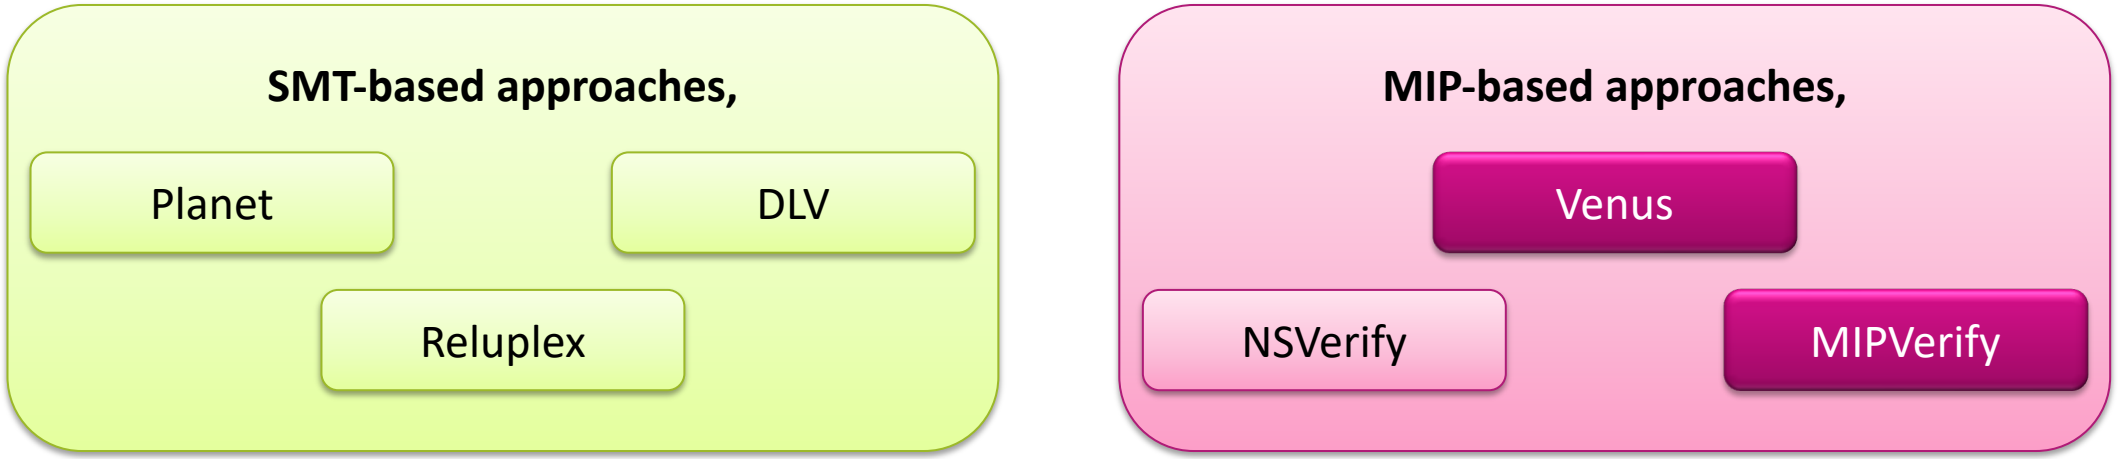

General workflow of a MIP-based verifier

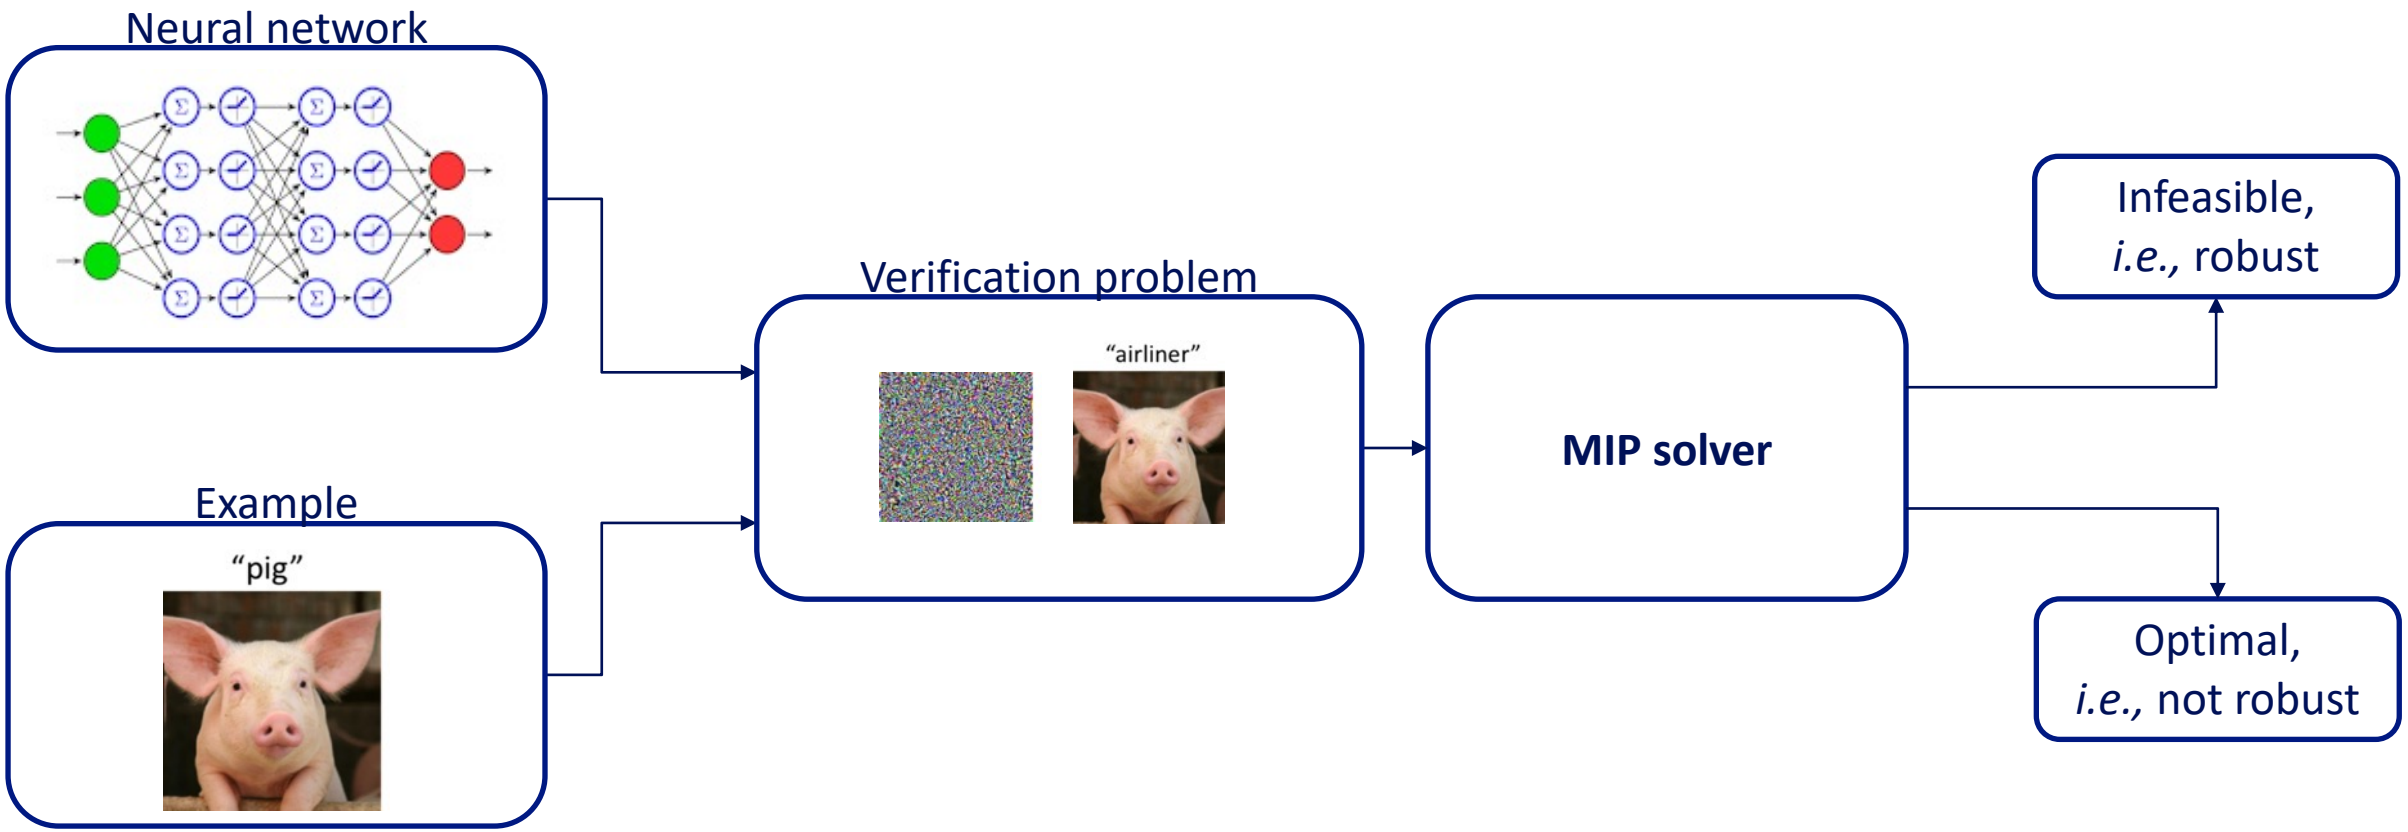

General workflow of a MIP-based verifier

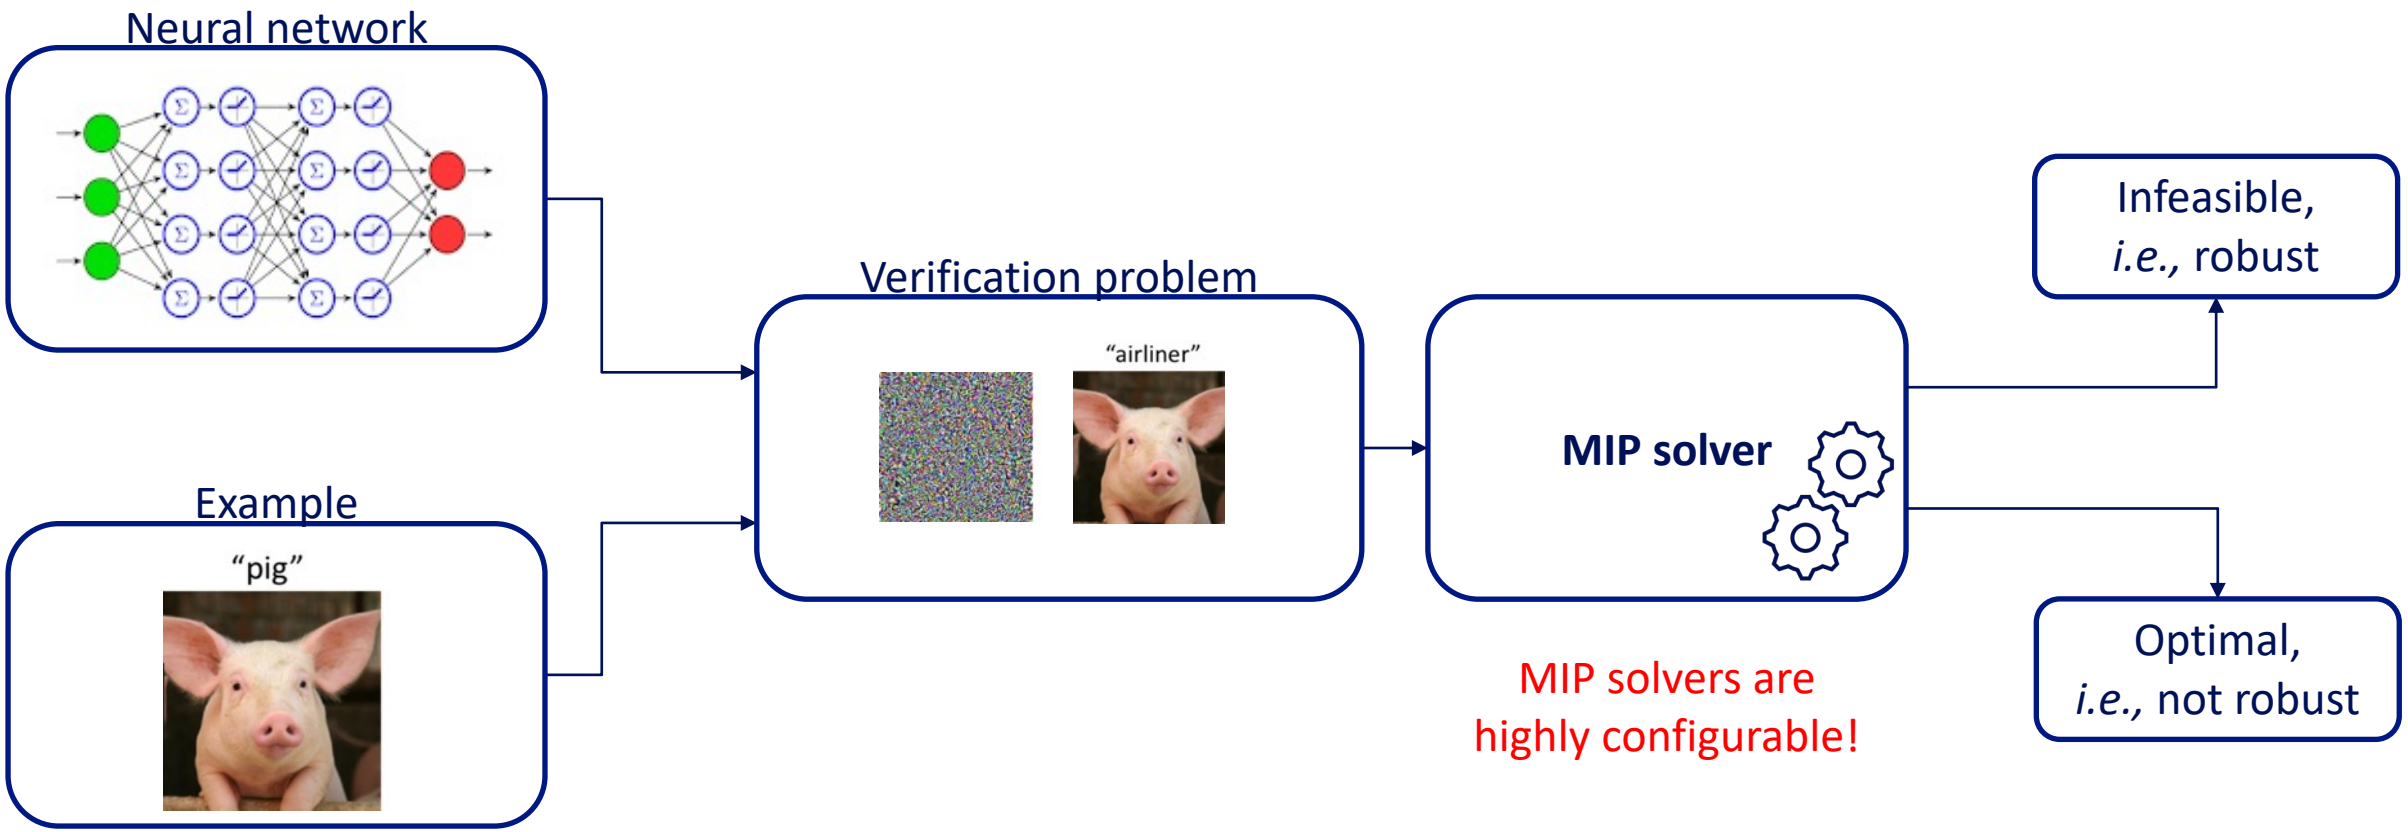

Main idea: Automated configuration of MIP solvers

- Only succeeds if instance set is homogenous

Main idea: Automated configuration of MIP solvers

- Only succeeds if instance set is homogenous

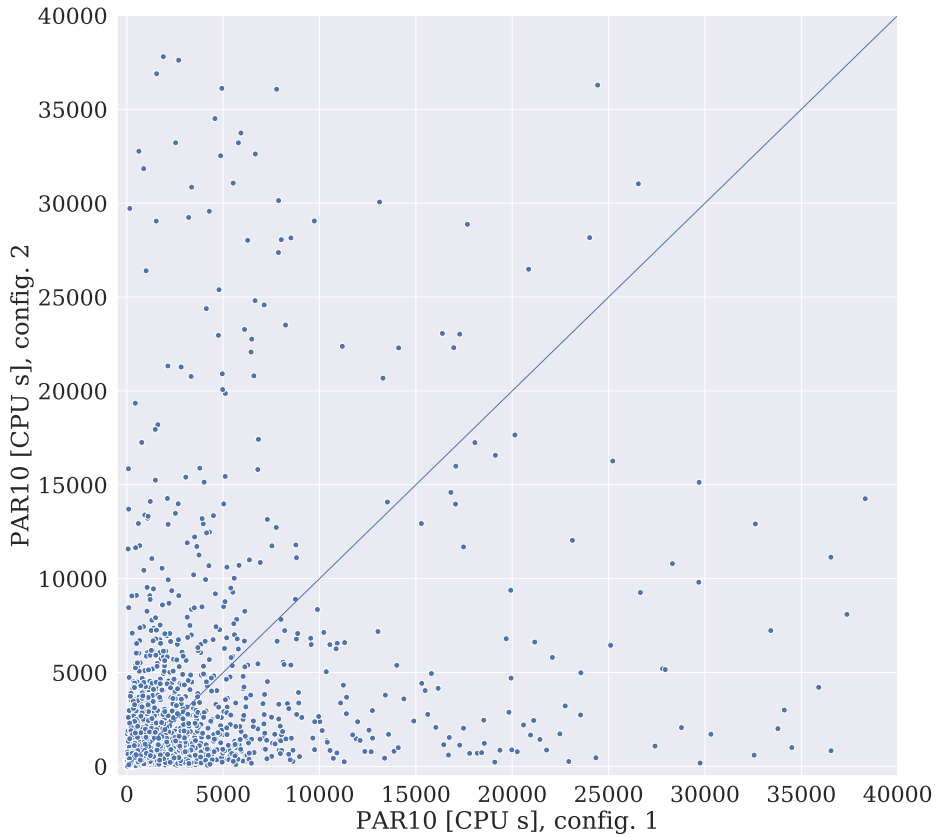

Workflow of our proposed solution

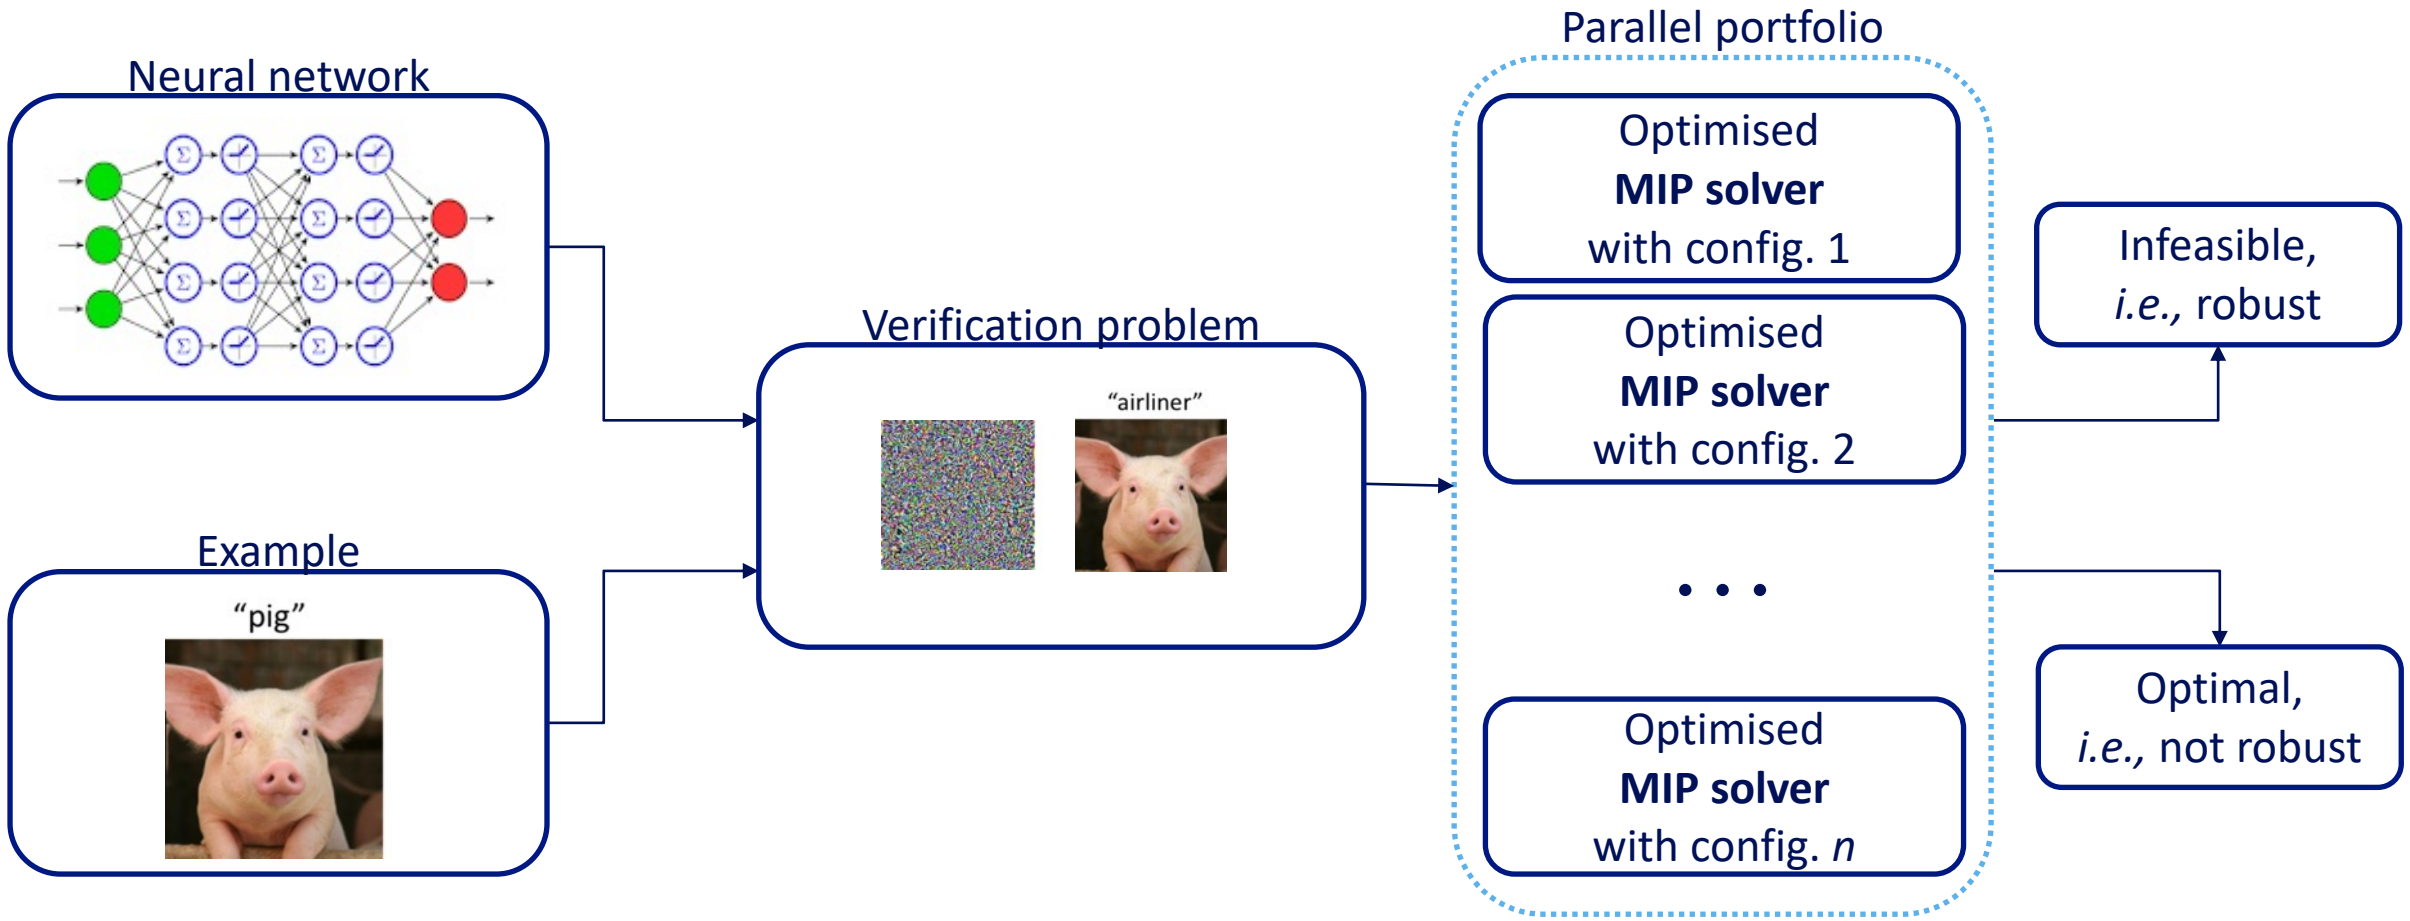

Our approach outperforms state-of-the-art approaches

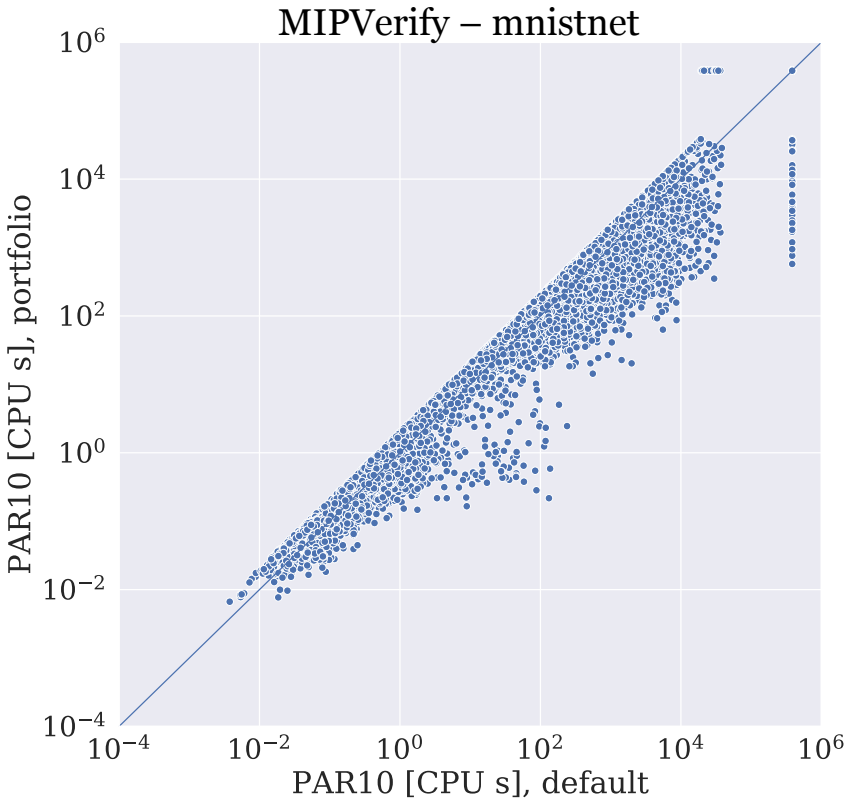

→ 1.61–fold speedup

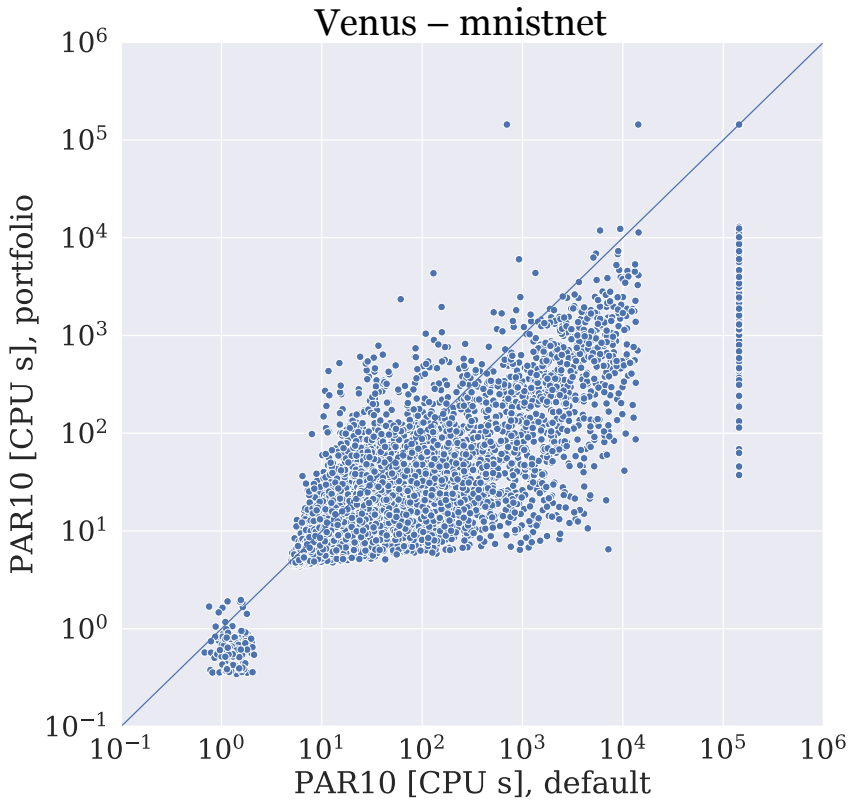

→ 7.26–fold speedup

Our approach outperforms state-of-the-art approaches

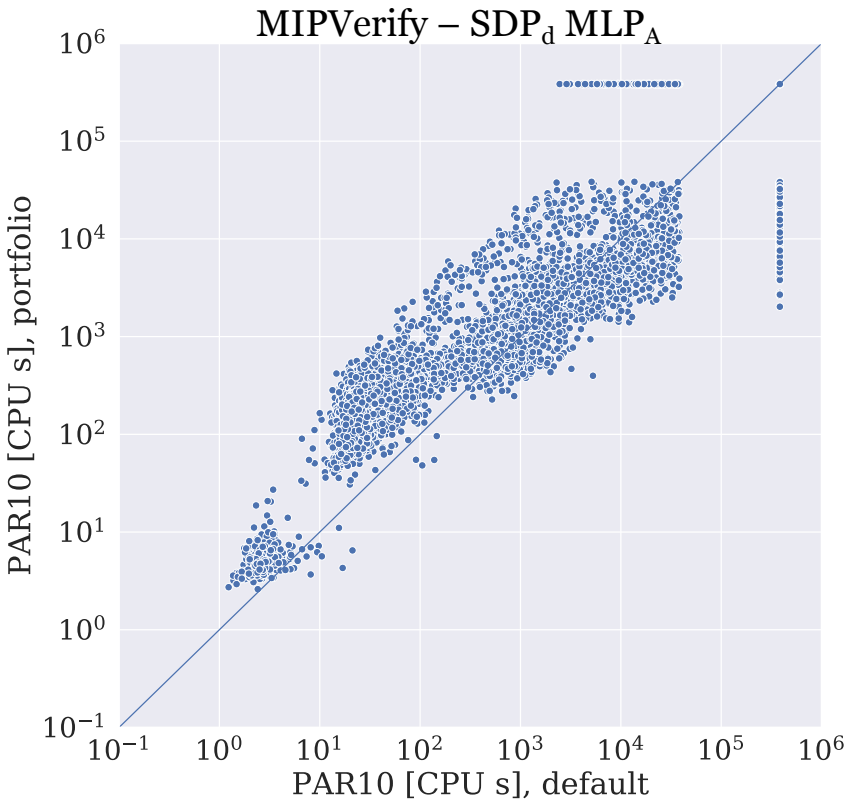

→ 4.7–fold speedup

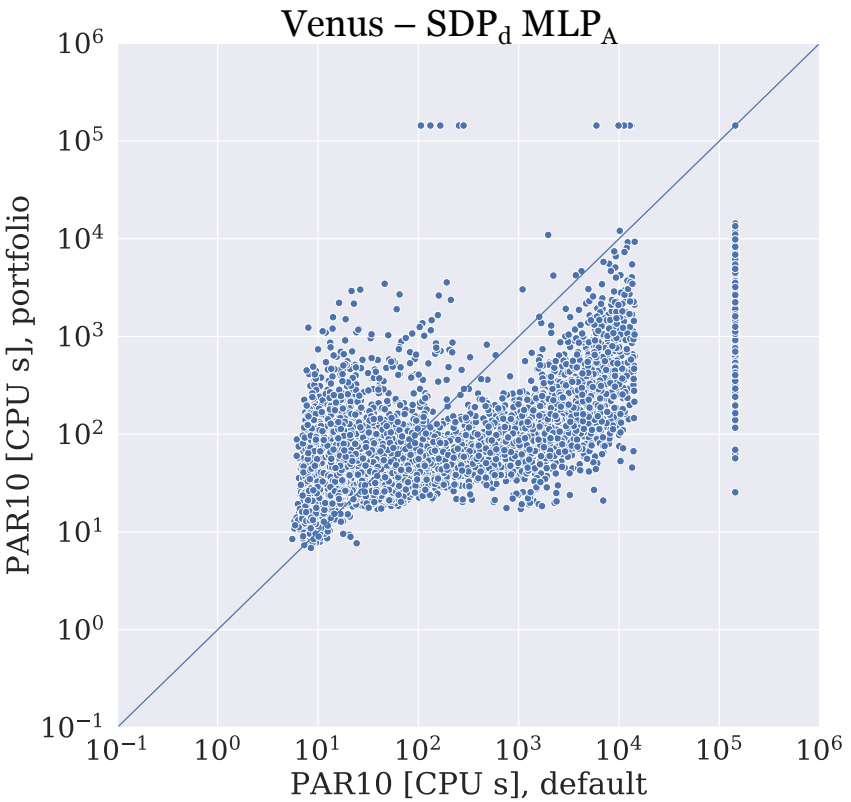

→ 10.3–fold speedup

Conclusions

- **Automated algorithm configuration and portfolio construction** techniques can strongly improve the performance of neural network verification algorithms
- More specifically, we achieved substantial improvements over SOTA methods employed at default, in terms of CPU running time, timeouts and adversarial error bounds
- Future work involves automated selection and extension to further hyperparameters

[König, Hoos, van Rijn. Speeding Up Neural Network Robustness Verification via Algorithm Configuration and an Optimised Mixed Integer Linear Programming Solver Portfolio. *Machine Learning.* 2022.]

Does the observed heterogeneity of MIP-encoded verification problem instances generalise to other types of verification problem instances?

Critical assessment of neural network verifiers

[König, Bosman, Hoos, van Rijn. Critically Assessing the State of the Art in CPU-based Local Robustness Verification. *Workshop on Artificial Intelligence Safety @AAAI*. 2023]

Critical assessment of neural network verifiers

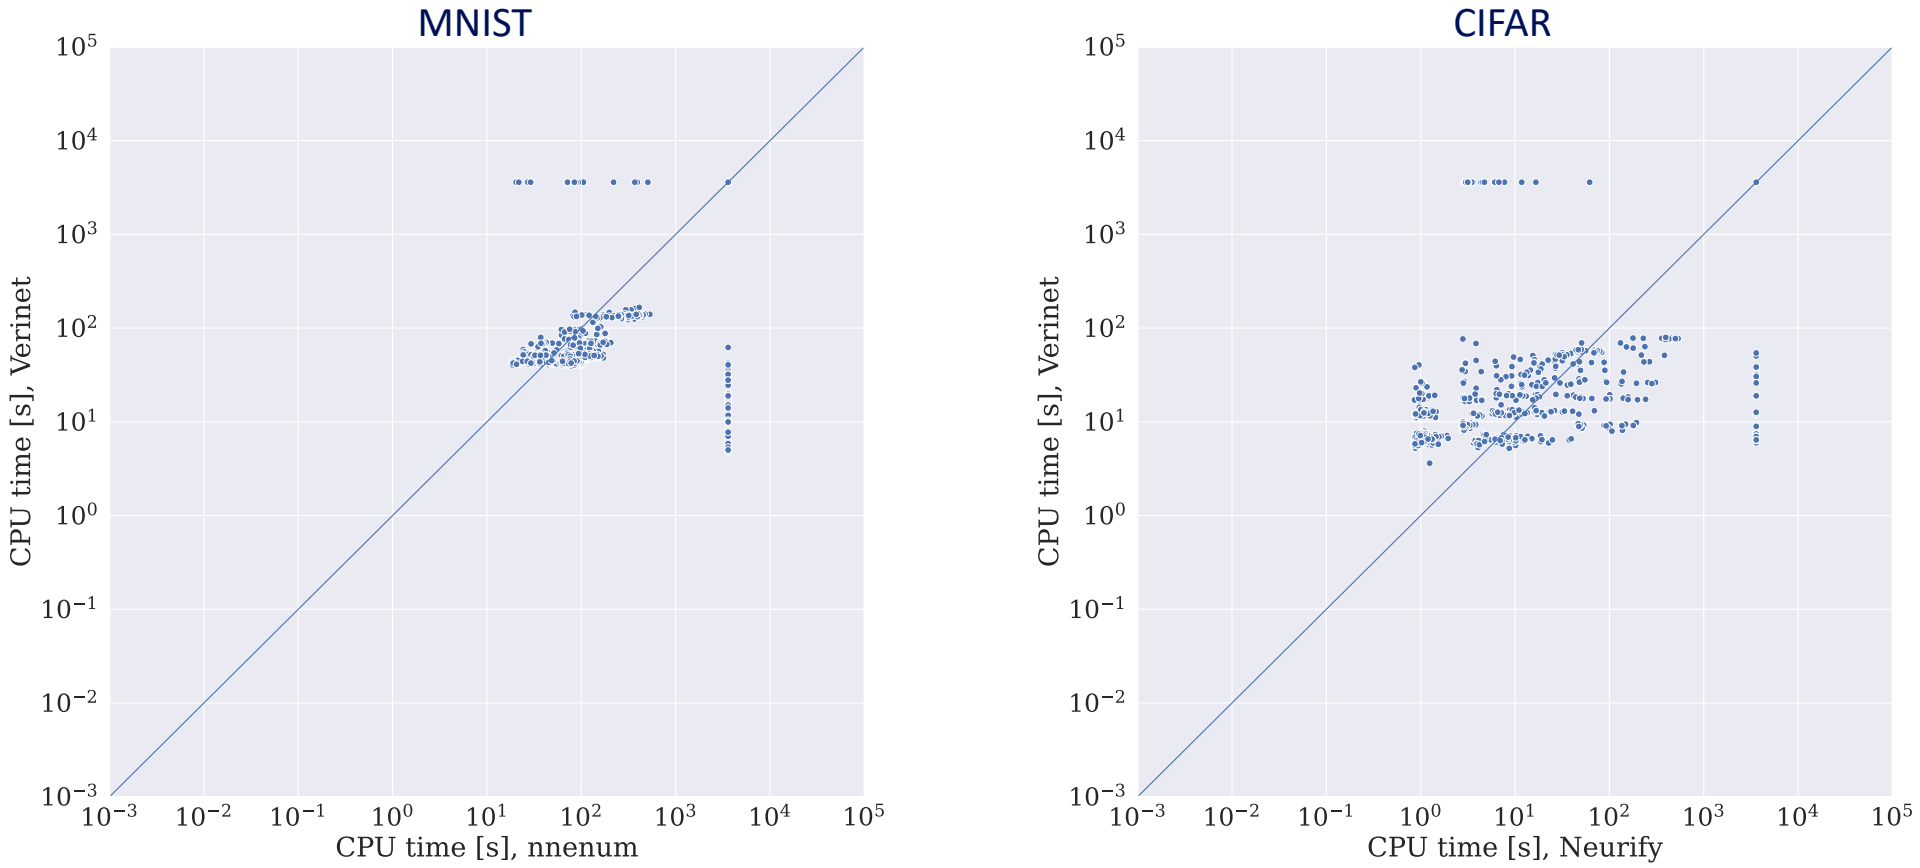

[König, Bosman, Hoos, van Rijn. Critically Assessing the State of the Art in CPU-based Local Robustness Verification. *Workshop on Artificial Intelligence Safety @AAAI*. 2023]

Vision: Auto-Verify for neural network verification